

<https://www.kaggle.com/code/moulabhaskar/pytorch-rice-classification/data>

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

import matplotlib.pyplot as plt
import numpy as np
import pathlib
import os

import copy
import random
import time

/usr/local/lib/python3.9/site-packages/tqdm/auto.py:22: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm

SEED = 1234
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

path = '/Users/cassidy/rice-classification/Rice_Image_Dataset/new-
dataset'

data_dir = pathlib.Path(path)
data_dir

PosixPath('/Users/cassidy/rice-classification/Rice_Image_Dataset/new-
dataset')

arborio = list(data_dir.glob('Arborio/*'))[:1]
basmati = list(data_dir.glob('Basmati/*'))[:1]
ipsala = list(data_dir.glob('Ipsala/*'))[:1]
jasmine = list(data_dir.glob('Jasmine/*'))[:1]
karacadag = list(data_dir.glob('Karacadag/*'))[:1]

import matplotlib.image as img
fig, ax = plt.subplots(ncols = 5, figsize = (20,5))
fig.suptitle('Rice Category')
arborio_image = img.imread(arborio[0])
```

```

basmati_image = img.imread(basmati[0])
ipsala_image = img.imread(ipsala[0])
jasmine_image = img.imread(jasmine[0])
karacadag_image = img.imread(karacadag[0])

```

```

ax[0].set_title('arborio')
ax[1].set_title('basmati')
ax[2].set_title('ipsala')
ax[3].set_title('jasmine')
ax[4].set_title('karacadag')

```

```

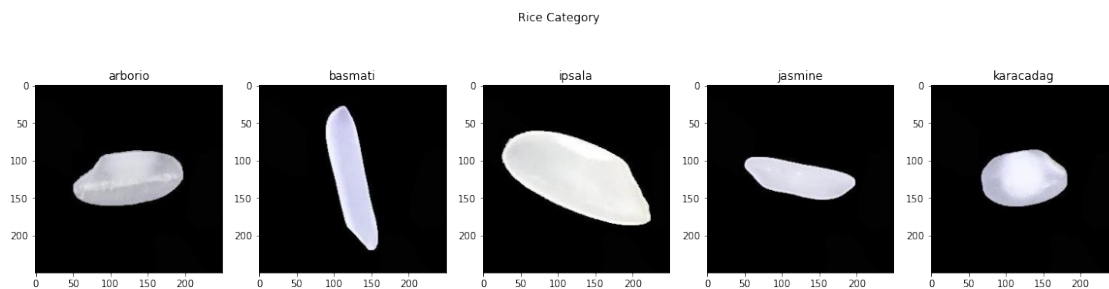
ax[0].imshow(arborio_image)
ax[1].imshow(basmati_image)
ax[2].imshow(ipsala_image)
ax[3].imshow(jasmine_image)
ax[4].imshow(karacadag_image)

```

```

plt.show()

```



```

df_labels = {
    'Arborio' : 0,
    'Basmati' : 1,
    'Ipsala' : 2,
    'Jasmine' : 3,
    'Karacadag': 4
}

```

```

class_names = ['Karacadag', 'Basmati', 'Jasmine', 'Arborio', 'Ipsala']
num_class = len(class_names)

```

```

image_files = [[os.path.join(path, class_name, x) for x in
os.listdir(os.path.join(path, class_name))] for class_name in
class_names]

```

```

images_paths = []
for i in range(5):
    for j in range(len(image_files[i])):
        current = image_files[i]
        images_paths.append(current[j])

```

```

import random

random.shuffle(images_paths)

train_paths = images_paths[:2100]
test_paths = images_paths[2100:2350]
valid_paths = images_paths[2350:]

import torchvision.transforms as transforms

transformations = transforms.Compose(
    [
        transforms.Resize((256,256)),
        transforms.RandomHorizontalFlip(),
        transforms.RandomRotation(30),
        transforms.ToTensor(),
        transforms.Normalize((0.5),(0.5))
    ]
)

from torch.utils.data import Dataset
from random import randint
from PIL import Image
class ImageDataset(Dataset):
    def __init__(self, df_labels, base_dir, transform=None):
        super().__init__()

        self.base_dir = base_dir
        self.df_labels = df_labels
        self.transform = transform

    def __len__(self):
        return len(self.base_dir)

    def __getitem__(self, index):

        image_path = self.base_dir[index]

        image = Image.open(image_path)

        label_name = image_path.split('/')[-2]

        label = self.df_labels[label_name]

        if self.transform is not None:
            image = self.transform(image)
        return (image, label)

train_data = ImageDataset(df_labels, train_paths, transformations)
test_data = ImageDataset(df_labels, test_paths, transformations)
valid_data = ImageDataset(df_labels, valid_paths, transformations)

```

```

from torch.utils.data import DataLoader
train_ds = DataLoader(train_data, batch_size = 64, shuffle = True)
val_ds = DataLoader(valid_data, batch_size = 64, shuffle = True)
test_ds = DataLoader(test_data, batch_size = 64, shuffle = True)

class LeNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels=3,
                                out_channels=6,
                                kernel_size=5)

        self.conv2 = nn.Conv2d(in_channels=6,
                                out_channels=16,
                                kernel_size=5)

        self.fc_1 = nn.Linear(16 * 61 * 61, 120)
        self.fc_2 = nn.Linear(120, 84)
        self.fc_3 = nn.Linear(84, output_dim)

    def forward(self, x):
        #(3, 256, 256) ---> input
        x = self.conv1(x)
        #(6, 252, 252) ---> output
        x = F.max_pool2d(x, kernel_size=2)
        #(6, 126, 126)
        x = F.relu(x)
        x = self.conv2(x)
        #(16, 122, 122)
        x = F.max_pool2d(x, kernel_size=2)
        #(16, 61, 61)
        x = F.relu(x)
        x = x.view(x.shape[0], -1)
        h = x
        x = self.fc_1(x)
        x = F.relu(x)
        x = self.fc_2(x)
        x = F.relu(x)
        x = self.fc_3(x)

        return x, h

OUTPUT_DIM = 5
model = LeNet(OUTPUT_DIM)

def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if
                p.requires_grad)

```

```
print(f'The model has {count_parameters(model):,} trainable  
parameters')
```

The model has 7,157,901 trainable parameters

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
model = model.to(device)
```

```
from torch.optim import Adam
```

```
optimizer = Adam(model.parameters(), lr = 0.001)
```

```
loss_fun = F.cross_entropy
```

```
loss_train = []
```

```
acc_train = []
```

```
loss_val = []
```

```
acc_val = []
```

```
def train(model, epoch, train_ds):
```

```
    model.train()
```

```
    total_num = len(train_ds.dataset)
```

```
    train_loss = 0
```

```
    correct_num = 0
```

```
    for image, label in train_ds:
```

```
        image = image.to(device)
```

```
        label = label.to(device)
```

```
        # Convert the tag from int32 type to long type, otherwise the  
calculation loss will report an error
```

```
        label = label.to(torch.long)
```

```
        output, _ = model(image)
```

```
        loss = loss_fun(output, label)
```

```
        train_loss += loss.item() * label.size(0)
```

```
        optimizer.zero_grad()
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
        predict = torch.argmax(output, dim=-1)
```

```
        correct_num += label.eq(predict).sum()
```

```
    train_loss = train_loss / total_num
```

```
    train_acc = correct_num / total_num
```

```
    print('epoch: {} --> train_loss: {:.6f} - train_acc: {:.6f} -  
' .format(
```

```
        epoch, train_loss, train_acc), end='')
```

```
    loss_train.append(train_loss)
```

```
acc_train.append(train_acc)
```

```
def evaluate(model, eval_ds, mode='val'):  
    model.eval()  
  
    total_num = len(eval_ds.dataset)  
    eval_loss = 0  
    correct_num = 0  
  
    for image, label in eval_ds:  
        image = image.to(device)  
        label = label.to(device)  
        label = label.to(torch.long)  
  
        output, _ = model(image)  
        loss = loss_fun(output, label)  
        eval_loss += loss.item() * label.size(0)  
  
        predict = torch.argmax(output, dim=-1)  
        correct_num += label.eq(predict).sum()  
  
    eval_loss = eval_loss / total_num  
    eval_acc = correct_num / total_num  
  
    print('{}_loss: {:.6f} - {}_acc: {:.6f}'.format(  
        mode, eval_loss, mode, eval_acc))  
  
    loss_val.append(eval_loss)  
    acc_val.append(eval_acc)
```

```
for epoch in range(5):  
    train(model, epoch, train_ds)  
    evaluate(model, val_ds)
```

```
epoch: 0 --> train_loss: 0.348852 - train_acc: 0.889524 - val_loss:  
0.385942 - val_acc: 0.800000  
epoch: 1 --> train_loss: 0.171093 - train_acc: 0.936667 - val_loss:  
0.098183 - val_acc: 0.960000  
epoch: 2 --> train_loss: 0.121526 - train_acc: 0.955238 - val_loss:  
0.077174 - val_acc: 0.966667  
epoch: 3 --> train_loss: 0.083975 - train_acc: 0.971429 - val_loss:  
0.024736 - val_acc: 0.993333  
epoch: 4 --> train_loss: 0.056167 - train_acc: 0.978571 - val_loss:  
0.061933 - val_acc: 0.973333
```

```
pred = []  
actual = []
```

```
def test(model, eval_ds, mode='val'):
```

```

model.eval()

total_num = len(eval_ds.dataset)
eval_loss = 0
correct_num = 0

for image, label in eval_ds:
    image = image.to(device)
    label = label.to(device)
    label = label.to(torch.long)

    output, _ = model(image)
    loss = loss_fun(output, label)
    eval_loss += loss.item() * label.size(0)

    predict = torch.argmax(output, dim=-1)
    correct_num += label.eq(predict).sum()

    pred.append(predict.numpy())
    actual.append(label.numpy())

eval_loss = eval_loss / total_num
eval_acc = correct_num / total_num

print('{}_loss: {:.6f} - {}_acc: {:.6f}'.format(
    mode, eval_loss, mode, eval_acc))

loss_val.append(eval_loss)
acc_val.append(eval_acc)

return pred, actual

y_pred, y = test(model, test_ds, mode='test')
test_loss: 0.100187 - test_acc: 0.968000

loss_val.pop()
0.08778621548414231

acc_val.pop()
tensor(0.9600)


epochs = [0,1,2,3,4]


plt.plot(epochs, acc_train, 'r', label='Training acc')
plt.plot(epochs, acc_val, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')

```

```
plt.legend()
plt.figure()

plt.plot(epochs, loss_train, 'r', label='Training loss')
plt.plot(epochs, loss_val, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend()
plt.show()
```





```
import itertools
from sklearn.metrics import confusion_matrix

def draw_confusion_matrix(y, yhat, classes):
    """
    Draws a confusion matrix for the given target and predictions
    Adapted from scikit-learn and discussion example.
    """
    plt.cla()
    plt.clf()
    matrix = confusion_matrix(y, yhat)
    plt.imshow(matrix, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title("Confusion Matrix")
    plt.colorbar()
    num_classes = len(classes)
    plt.xticks(np.arange(num_classes), classes, rotation=90)
    plt.yticks(np.arange(num_classes), classes)

    fmt = 'd'
    thresh = matrix.max() / 2.
    for i, j in itertools.product(range(matrix.shape[0]),
    range(matrix.shape[1])):
        plt.text(j, i, format(matrix[i, j], fmt),
                horizontalalignment="center",
                color="white" if matrix[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```

plt.tight_layout()
plt.show()

import tensorflow as tf

y_new = []

for batch in y:
    for i in batch:
        y_new.append(batch[i])

y = y_new

y_pred_new = []

for batch in y_pred:
    for i in batch:
        y_pred_new.append(batch[i])

y_pred = y_pred_new

classes = ['Arborio', 'Basmati', 'Ipsala', 'Jasmine', 'Karacadag']

draw_confusion_matrix(y_new, y_pred_new, classes)

```

