

```

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # this is used for the plot the graph
import os
from sklearn.model_selection import train_test_split, cross_val_score,
GridSearchCV
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import confusion_matrix
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.cluster import KMeans
from sklearn.model_selection import KFold
from sklearn import model_selection

import itertools

import warnings
warnings.filterwarnings('ignore')

from google.colab import drive
drive.mount('drive')

def draw_confusion_matrix(y, yhat, classes):
    """
    Draws a confusion matrix for the given target and predictions
    Adapted from scikit-learn and discussion example.
    """
    plt.cla()
    plt.clf()
    matrix = confusion_matrix(y, yhat, labels=classes)
    plt.imshow(matrix, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title("Confusion Matrix")
    plt.colorbar()
    num_classes = len(classes)
    plt.xticks(np.arange(num_classes), classes, rotation=90)
    plt.yticks(np.arange(num_classes), classes)

    fmt = 'd'
    thresh = matrix.max() / 2.
    for i, j in itertools.product(range(matrix.shape[0]),
range(matrix.shape[1])):
        plt.text(j, i, format(matrix[i, j], fmt),
                horizontalalignment="center",
                color="white" if matrix[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')

```

```
plt.tight_layout()
plt.show()
```

```
rice_data = pd.read_csv("/content/drive/MyDrive/CSSI - Intermediate
Course - Data Science/Final Project/Rice_MSC_Dataset.csv")
rice_data.head()
```

Mounted at drive

	AREA	PERIMETER	MAJOR_AXIS	MINOR_AXIS	ECCENTRICITY	EQDIASQ
SOLIDITY \						
0	7805	437.915	209.8215	48.0221	0.9735	99.6877
0.9775						
1	7503	340.757	138.3361	69.8417	0.8632	97.7400
0.9660						
2	5124	314.617	141.9803	46.5784	0.9447	80.7718
0.9721						
3	7990	437.085	201.4386	51.2245	0.9671	100.8622
0.9659						
4	7433	342.893	140.3350	68.3927	0.8732	97.2830
0.9831						

	CONVEX_AREA	EXTENT	ASPECT_RATIO	...	ALLdaub4L	ALLdaub4a
ALLdaub4b \						
0	7985	0.3547	4.3693	...	113.9924	65.0610
59.5989						
1	7767	0.6637	1.9807	...	105.7055	64.3685
62.2084						
2	5271	0.4760	3.0482	...	109.7155	62.6423
58.7439						
3	8272	0.6274	3.9325	...	116.5405	64.9069
60.2562						
4	7561	0.6006	2.0519	...	107.7502	64.7071
61.3549						

	ALLdaub4Y	ALLdaub4Cb	ALLdaub4Cr	ALLdaub4XX	ALLdaub4YY
ALLdaub4ZZ \					
0	104.8552	67.8779	63.0828	0.3673	0.3793
0.4733					
1	96.8375	65.5371	63.5832	0.3014	0.3144
0.3641					
2	100.2352	68.9753	59.8342	0.3233	0.3445
0.4448					
3	107.2560	67.3298	63.2237	0.3880	0.4020
0.4904					
4	98.8704	66.2048	63.5378	0.3184	0.3303
0.3928					

CLASS

```

0 Basmati
1 Arborio
2 Jasmine
3 Basmati
4 Arborio

```

```
[5 rows x 107 columns]
```

```

#rice_data = rice_data.drop(columns=["CLASS"])
rice_data

```

	AREA	PERIMETER	MAJOR_AXIS	MINOR_AXIS	ECCENTRICITY	
EQDIASQ \						
0	7805	437.915	209.8215	48.0221	0.9735	
99.6877						
1	7503	340.757	138.3361	69.8417	0.8632	
97.7400						
2	5124	314.617	141.9803	46.5784	0.9447	
80.7718						
3	7990	437.085	201.4386	51.2245	0.9671	
100.8622						
4	7433	342.893	140.3350	68.3927	0.8732	
97.2830						
...	...	...	...	...	...	..
.						
74995	5551	285.911	114.1695	62.9079	0.8345	
84.0699						
74996	7696	322.703	121.3900	81.1375	0.7438	
98.9892						
74997	7579	339.295	136.3125	71.2866	0.8524	
98.2338						
74998	15174	489.502	200.9486	97.6282	0.8740	
138.9969						
74999	12931	452.635	185.5138	90.2651	0.8736	
128.3131						
	SOLIDITY	CONVEX_AREA	EXTENT	ASPECT_RATIO	...	ALLdaub4L
ALLdaub4a \						
0	0.9775	7985	0.3547	4.3693	...	113.9924
65.0610						
1	0.9660	7767	0.6637	1.9807	...	105.7055
64.3685						
2	0.9721	5271	0.4760	3.0482	...	109.7155
62.6423						
3	0.9659	8272	0.6274	3.9325	...	116.5405
64.9069						
4	0.9831	7561	0.6006	2.0519	...	107.7502
64.7071						
...	...	...	...	...	...	...
...						

74995	0.9846	5638	0.6418	1.8149	...	103.9529
64.9225						
74996	0.9868	7799	0.7309	1.4961	...	108.9778
65.4571						
74997	0.9805	7730	0.6399	1.9122	...	106.0881
64.1869						
74998	0.9766	15537	0.7903	2.0583	...	119.2037
63.3545						
74999	0.9760	13249	0.7640	2.0552	...	121.4198
63.5424						

	ALLdaub4b	ALLdaub4Y	ALLdaub4Cb	ALLdaub4Cr	ALLdaub4XX
ALLdaub4YY \					
0	59.5989	104.8552	67.8779	63.0828	0.3673
0.3793					
1	62.2084	96.8375	65.5371	63.5832	0.3014
0.3144					
2	58.7439	100.2352	68.9753	59.8342	0.3233
0.3445					
3	60.2562	107.2560	67.3298	63.2237	0.3880
0.4020					
4	61.3549	98.8704	66.2048	63.5378	0.3184
0.3303					
...	...	...	...	...	...
...					
74995	62.4355	95.2780	65.5114	64.4457	0.2895
0.2997					
74996	59.9502	100.2301	67.5089	63.6028	0.3335
0.3426					
74997	61.3876	97.1585	66.2445	63.0596	0.3028
0.3164					
74998	64.8200	109.3027	63.3122	63.5967	0.3970
0.4215					
74999	65.2355	111.4580	63.0129	63.9117	0.4162
0.4414					

	ALLdaub4ZZ	CLASS
0	0.4733	Basmati
1	0.3641	Arborio
2	0.4448	Jasmine
3	0.4904	Basmati
4	0.3928	Arborio
...	...	...
74995	0.3455	Arborio
74996	0.4257	Karacadag
74997	0.3761	Arborio
74998	0.4469	Ipsala
74999	0.4626	Ipsala

[75000 rows x 107 columns]

```
rice_data = rice_data.drop(rice_data.iloc[:, 16:106], axis=1)
rice_data.head()
```

	AREA SOLIDITY	PERIMETER \	MAJOR_AXIS	MINOR_AXIS	ECCENTRICITY	EQDIASQ
0	7805	437.915	209.8215	48.0221	0.9735	99.6877
	0.9775					
1	7503	340.757	138.3361	69.8417	0.8632	97.7400
	0.9660					
2	5124	314.617	141.9803	46.5784	0.9447	80.7718
	0.9721					
3	7990	437.085	201.4386	51.2245	0.9671	100.8622
	0.9659					
4	7433	342.893	140.3350	68.3927	0.8732	97.2830
	0.9831					

	CONVEX_AREA SHAPEFACTOR_1	EXTENT \	ASPECT_RATIO	ROUNDNESS	COMPACTNESS
0	7985	0.3547	4.3693	0.5114	0.4751
	0.0269				
1	7767	0.6637	1.9807	0.8120	0.7065
	0.0184				
2	5271	0.4760	3.0482	0.6505	0.5689
	0.0277				
3	8272	0.6274	3.9325	0.5256	0.5007
	0.0252				
4	7561	0.6006	2.0519	0.7944	0.6932
	0.0189				

	SHAPEFACTOR_2	SHAPEFACTOR_3	SHAPEFACTOR_4	CLASS
0	0.0062	0.2257	0.9863	Basmati
1	0.0093	0.4992	0.9888	Arborio
2	0.0091	0.3236	0.9865	Jasmine
3	0.0064	0.2507	0.9859	Basmati
4	0.0092	0.4806	0.9860	Arborio

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.base import BaseEstimator, TransformerMixin
```

```
rice_data_target = rice_data["CLASS"]
rice_data = rice_data.drop(["CLASS"], axis = 1)
```

```
numerical_features = list(rice_data)
```

```
pipeline = ColumnTransformer([
    ('std_scaler', StandardScaler(), numerical_features)
```

```

    ])

rice_data_prepared = pipeline.fit_transform(rice_data)

X_train, X_test, y_train, y_test =
train_test_split(rice_data_prepared, rice_data_target, test_size =
0.2)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(60000, 16) (15000, 16) (60000,) (15000,)

#Logistic Regression w/o Penalty

log_reg = LogisticRegression(multi_class="multinomial",
solver="lbfgs", penalty='none') #maybe make into knn or svm
log_reg.fit(X_train, y_train)


y_pred = log_reg.predict(X_test)
log_score = log_reg.predict_proba(X_test)

print("%-12s %f" % ('Accuracy:',
metrics.accuracy_score(y_test,y_pred)))
print("%-12s %f" % ('Precision:',
metrics.precision_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('Recall:',
metrics.recall_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('F1 Score:',
metrics.f1_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))

Accuracy:      0.977000
Precision:     0.977167
Recall:        0.977190
F1 Score:      0.977166

draw_confusion_matrix(y_test, y_pred, ['Arborio', 'Basmati', 'Ipsala',
'Jasmine', 'Karacadag'])
#rice_data.head(10)

```



#### *#Logistic Regression w/ l2 Penalty*

```
log_reg = LogisticRegression(multi_class="multinomial",
                             solver="lbfgs", penalty='l2')
log_reg.fit(X_train, y_train)


y_pred = log_reg.predict(X_test)
log_score = log_reg.predict_proba(X_test)

print("%-12s %f" % ('Accuracy:',
metrics.accuracy_score(y_test,y_pred)))
print("%-12s %f" % ('Precision:',
metrics.precision_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('Recall:',
metrics.recall_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('F1 Score:',
metrics.f1_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))

draw_confusion_matrix(y_test, y_pred, ['Arborio', 'Basmati', 'Ipsala',
'Jasmine', 'Karacadag'])

Accuracy:    0.976267
Precision:   0.976459
```

Recall: 0.976441  
F1 Score: 0.976437



```
'''
#KNN Optimizer
k_values = [1,2,3,5,7,9,10,20,200]
for k in k_values:
    knnclass = KNeighborsClassifier(n_neighbors=k)
    knnclass.fit(X_train, y_train)
    y_pred = knnclass.predict(X_test)
    print(k, "%-12s %f" % ('Accuracy:',
metrics.accuracy_score(y_test,y_pred)))
'''

{"type": "string"}

#KNN

knnclass = KNeighborsClassifier(n_neighbors=5)
knnclass.fit(X_train, y_train)
y_pred = knnclass.predict(X_test)

print("%-12s %f" % ('Accuracy:',
metrics.accuracy_score(y_test,y_pred)))
print("%-12s %f" % ('Precision:',
metrics.precision_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('Recall:',
```



```

metrics.recall_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))
print("%-12s %f" % ('F1 Score:',
metrics.f1_score(y_test,y_pred,labels=None, pos_label=1,
average='macro', sample_weight=None)))


draw_confusion_matrix(y_test, y_pred, ['Arborio', 'Basmati', 'Ipsala',
'Jasmine', 'Karacadag'])

```

```

Accuracy:    0.978000
Precision:   0.978158
Recall:      0.978201
F1 Score:    0.978158

```



*#K Fold of Top 2*

```

kfold = model_selection.KFold(n_splits=10, random_state=42,
shuffle=True)

```

*# replace the two models below with what you picked*

```

knn_model_kfold = KNeighborsClassifier(n_neighbors=5)
log_model_kfold = LogisticRegression(multi_class="multinomial",
solver="lbfgs", penalty='none')

```

*# Finally we pull it all together. We call cross val score to generate an accuracy performance score for our model*

*# we define our learning model, data, labels, and cross-val splitting strategy (all defined previously)*

```
knn_results_kfold = model_selection.cross_val_score(knn_model_kfold,
X_train, y_train, cv=kfold)
log_results_kfold = model_selection.cross_val_score(log_model_kfold,
X_train, y_train, cv=kfold)

# Because we're collecting results from all runs, we take the mean
value
print("For KNN Classifier our mean accuracy across folds is: %.2f%%" %
(knn_results_kfold.mean()*100.0))
print("For Logistic Regression our mean accuracy across folds is: %.2f
%%" % (log_results_kfold.mean()*100.0))
```

For KNN Classifier our mean accuracy across folds is: 97.92%  
For Logistic Regression our mean accuracy across folds is: 97.76%